

Exploiting HBM for CNN Acceleration on FPGAs

Multiple Kernel Execution using PYNQ-for-XRT

Akshay Kamath, Callie Hao, SharC Lab, School of Electrical and Computer Engineering, Georgia Tech

Motivation

FPGAs are ideal for energy-efficient CNN acceleration.

Accelerating large-scale models is highly challenging.

- E.g. ResNet-50 with HD (720 x 1280) inputs
- Off-chip memory access is the primary bottleneck!

Typical approach to acceleration is via tiling of inputs.

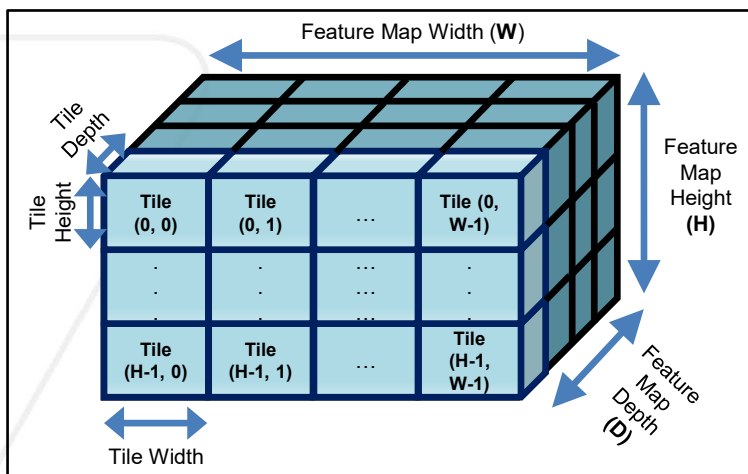
- Intra-tile parallelization limited by on-chip buffers
- Inter-tile parallelization limited by off-chip memory

HBM-based FPGAs have independent memory banks.

- Facilitates multiple tile-kernel execution in parallel!

PYNQ-for-XRT framework allows quick prototyping.

- Python-based host programming over OpenCL



Tiling-based Convolution for Acceleration of CNN Inference on FPGA

Tile Parallelization

Intra-tile parallelization involves optimizing convolution computation.

Inter-tile parallelization is achieved through multiple kernel instantiation.

- Each kernel is connected to different HBM memory banks

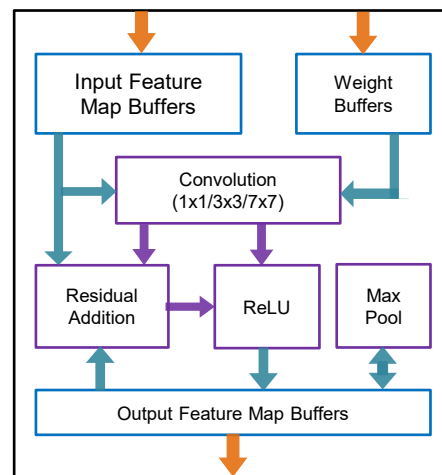
for each set of N tiles:

```
Sync input buffers from host to device memory
Start execution of all kernels in parallel
Wait for all kernels to finish execution
Sync output buffers from device memory to host
```


**Host
program
pseudo
-code**

```
Load input feature map tile from off-chip memory
Load weights and bias from off-chip memory
Perform convolution on-chip w/o ReLU
Store output feature map tile to off-chip memory
```

**Kernel
(device)
pseudo
-code**



HLS Convolution Tile Kernel Architecture



System-architecture for multiple kernel execution

Experimental Results

Demonstrated on CRNCH Lab's Alveo U50 FPGA

- Includes x86 CPU connected to 32 HBM banks
- Kernel developed using Xilinx Vitis HLS

**Setup for a sample ResNet50 3x3
convolution layer with BDD100K dataset**

IO Tile Dimensions	64 * 23 * 20
Weight Blocks Dimensions	64 * 64 * 3 * 3
Fixed-point Precision	32 bits
No. of Kernel Instances	4

Configuration	Avg. Latency (ms)
Single kernel	5.959
4 kernels (sequential)	23.935
4 kernels (parallel)	6.857

Parallel execution, thus, results in **3.5x** speedup!

- Less than ideal 4x speedup due to overhead of synchronization of data for 4 kernels