COSPlay: Using Coroutines for High-Throughput Synchronous Persistence

Marina Vemmou, Alexandros Daglis

School of Computer Science, Georgia Institute of Technology

Georgia Tech | Center for Research into Novel Computing Hierarchies

Challenges of Persistent Memory

- PM applications: need to be crash consistent
  - **Challenge: Volatile order ≠ persistent order**
- Must explicitly control order of PM updates
  - x86 persistency model

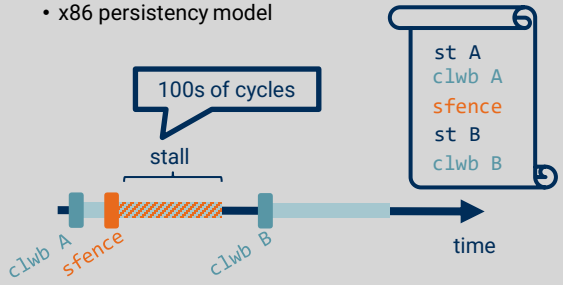

Persist Latency and BMOs

Backend Memory Operations (BMOs):

- performed before reaching persistent domain
- can add 100s of cycles on path to PM

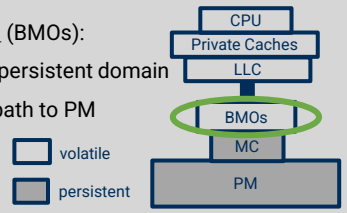

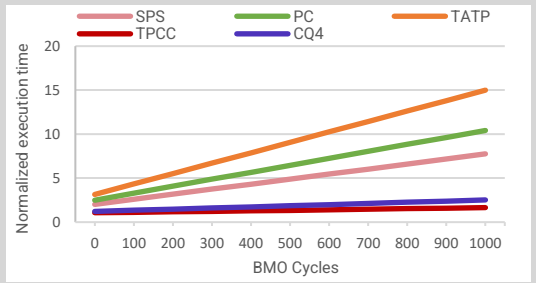

up to 15x

x86 persistence + high persist latency = slowdown up to 15x

Hiding Persist Latency Through Parallelism

- Key insight: many PM programs have inherent task parallelism

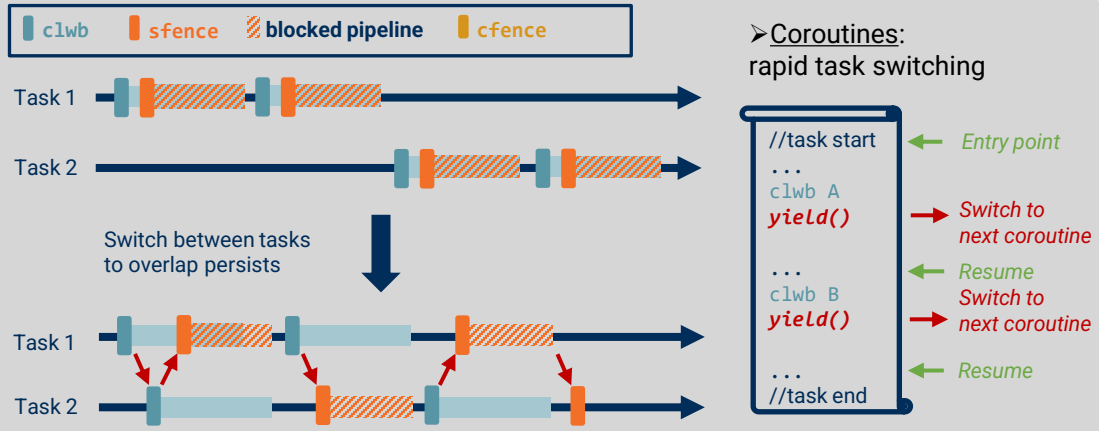

- Coroutines: rapid task switching

Coroutines for ns-scale task switching

Context-Aware Hardware for Selective Persist Ordering

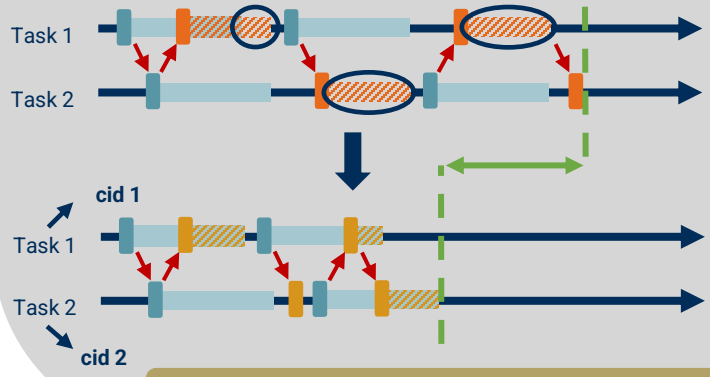

`sfence` waits for pending `clwb`s of all tasks

Insight: only need to maintain intra-task order

- Each coroutine given a context-id (cid)
- Each `clwb` associated with current cid
- Introducing `cfence`: waits only for `clwb`s of a *specific* cid

COSPlay: relax inter-task persist order, preserve intra-task persist order

Putting It All Together

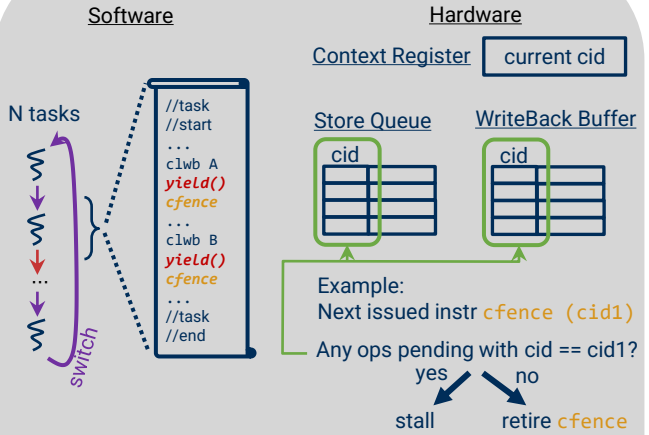

COSPlay Performance

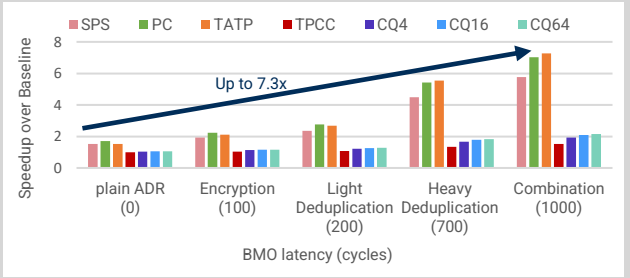

- Speedup even without BMOs
- Speedup increases with BMO latency

COSPlay improves throughput of PM programs by up to 7.3x

This work appeared at the 54th International Symposium on Microarchitecture (MICRO'21)